

AI meets Real-Time Bidding: A Case Study

Jian Xu

Big Data and Artificial Intelligence Lab TouchPal Inc.

July 12, 2017



Online Advertising Landscape

Real-Time Bidding Mechanism

Case Study: AI Enables Better Bidding Strategies



Online Advertising Landscape

Real-Time Bidding Mechanism

Case Study: AI Enables Better Bidding Strategies



Online Advertising Landscape

Real-Time Bidding Mechanism

Case Study: AI Enables Better Bidding Strategies

Online Advertising Landscape



The compound annual revenue growth rate over the past ten years for online advertising is 16%. Mobile accounting for 50.52% of all revenue in 2016, surpassing non-mobile revenue for the first time.

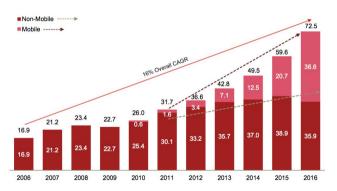


Figure: US online advertising annual revenue.

¹ Source: IAB/PwC Internet Ad Revenue Report, FY 2016



Online Advertising Landscape



Search ads (46%) and display ads (45%) contribute the majority of the online ad revenue ²



Figure: An search ad





Figure: A "native" display ad Figure: A traditional display ad

Xu et al. DACON'17 (TouchPal Inc.)

Visit of the second o

² Source: IAB/PwC Internet Ad Revenue Report, FY 2016

Online Advertising Landscape



- 80% US online display ad revenue comes from programmatic buying.
- 50% of the programmatic ad revenue comes from Real-Time Bidding (RTB) 3.

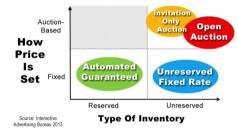


Figure: How display ad opportunities (inventory) are sold via programmatic

RTB is deeply powered by AI.

³ Source: https://www.emarketer.com/Article/eMarketer-Releases-New-Programmatic-Advertising-Estimates/1015682



Online Advertising Landscape

Real-Time Bidding Mechanism

Case Study: AI Enables Better Bidding Strategies

Real-Time Bidding Mechanism



Advertising buyers bid on an impression and, if the bid is won, the buyer's ad is instantly displayed on the publisher's site.

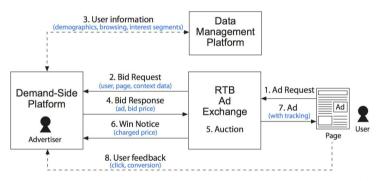


Figure: An illustration of how an ad impression is sold in real-time bidding.

⁴ Source: Weinan Zhang, Optimal Real-Time Bidding for Display Advertising, PhD Thesis, * 2016. * 🚊 🕨 🔻 🚊 🔻 🗢 🗸 🔾

Real-Time Bidding Mechanism



From the buyers' (advertisers, DSPs) perspective, it is important to bid smartly. There are two widely adopted problem settings:

No budget limit, maximize ROI

Suppose the return of an action (conversion) is CPA, the estimated action rate is \hat{AR} . The optimal bidding strategy is surprisingly simple

$$b_{OPT} = CPA \cdot \hat{AR} \tag{1}$$

With budget and lifetime limit, maximize some KPI

To optimize total action number within budget B and bid request volume T, Suppose p and F are p.d.f. and C.D.F. of the action rate respectively and the estimated action rate is $A\hat{R}$. The optimal bidding strategy is g

$$b_{OPT} = \frac{B}{T_{\theta 0}^{3} \vartheta(n-1)F(\vartheta)^{n-2}p(\vartheta)d\vartheta p(\theta)d\theta} \cdot \hat{A}R$$
 (2)

^a Source: Weinan Zhang, Optimal Real-Time Bidding for Display Advertising, PhD Thesis, 2016.

Real-Time Bidding Mechanism



RTB has to be accomplished with the help of AI:

- <100ms latency requirement,
- >1M bid requests per second (QPS)

All implements the bidding strategies by estimating \hat{AR} in real-time; and even more ...

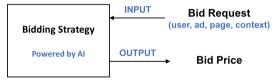


Figure: Real-Time Bidding as an AI task



Online Advertising Landscape

Real-Time Bidding Mechanism

Case Study: AI Enables Better Bidding Strategies

Case Study: AI Enables Better Bidding Strategy



The effective impact of an ad (AR lift) is more important than the absolute AR value.

A tiny example

Two users: a and b, campaign CPA: \$100

 AR_a : 0.04 if exposed to the ad, 0.03 if not (lift: 0.01); AR_b : 0.02 if exposed to the ad, 0.001 if not (lift: 0.019).

Bidder₁ bidsprop. to ARassuming exposed: \$4 for a, \$2 for b;

Bidder₂ bidsprop. to AR lift: \$2 for a, \$3.8 for b. Incremental value from

Bidder 1: 0.01 conversions;

Incremental value from Bidder 2: 0.019 conversions.

- If only one of them can be exposed to the ad, who will you select?
- Prevalent bidding strategy does not optimize campaign performance.

Predicting AR lift



Let ad be an ad, s be the state of a user at ad request time, and $s_+(ad)$ be the state of the user if ad is shown. Let p(action/s) be the AR of the user if ad is not shown and $p(action/s_+(ad))$ be the AR if ad is shown, the AR lift is

$$\Delta p = p(action/s+(ad)) - p(action/s)$$
 (3)

How to learn from data and predict Δp ?

Learn an *omnipotent* model to predict AR given any arbitrary state.

Use a function F to map a state to a set of features;

A generic AR prediction model \hat{P} is built upon the derived feature set; Finally the AR lift can be estimated as

$$\Delta \hat{p} = P(action | F(s+(ad))) - P(action | F(s))$$
 (4)



Predicting AR lift (cont.)



Traditional AR prediction models are trained based on impression/click events. Some concerns about the traditional models ⁵:

Not generalized enough to learn p(action/s), Suffering from survival bias,

Not leveraging all the action information.

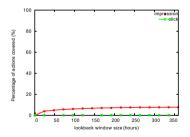


Figure: Only less than 10% of the converted users had been exposed to the ad of the advertiser.

Predicting AR lift (cont.)



Our approach to build a machine learning model to predict AR lift:

Training sample generation

Generate training samples on user + time-stamp level, Leverage all the action data (conversion pixel firings).

Features

```
\ellFrequency, Recency \ell \ell fimpression, click, conversion \ell \ell Frequency, Recency \ell \ell Search, page view \ell \ell Topic \ell \ell Topic \ell Demographics, geolocation, device, etc.
```

Model fitting

GBDT for rank order, Isotonic regression for calibration.

Online A/B test



Three equal-sized random buckets: No show: does

not bid at all,

Value-based bidding (50% budget assigned),

Lift-based bidding (50% budget assigned).

No show vs. value-based bidding

Adv	No bid		Value-base d bidding		Incremental action	Action lift
	# imps	# actions	#imps	# actions		
1	0	642	53,396	714	72	11.2%
2	0	823	298,333	896	73	8.9%
3	0	1,438	11,048,583	1,477	39	2.7%
4	0	1892	3,915,792	2,016	124	6.6%
5	0	5,610	6,015,322	6,708	1,098	19.6%

Table: Blind A/B test on five pilot advertisers - Value-based bidding v.s. "No bid".

Online A/B test (cont.)



No show vs. lift-based bidding

Adv	No bid		Lift-based bidding		Incremental action	Action lift
	# imps	# actions	#imps	# actions		
1	0	642	59,703	826	184	28.7%
2	0	823	431,637	980	157	19.1%
3	0	1,438	11,483,360	1509	71	4.9%
4	0	1892	4,368,441	2,471	579	30.6%
5	0	5,610	8,770,935	8,291	2,681	47.8%

Table: Blind A/B test on five pilot advertisers - Lift-based bidding v.s. "No bid".

Value-based bidding vs. lift-based bidding - Advertiser's perspective

Adv	Va	Value-based bidding			Lift-based bidding			Lift-over-lift
	#imps	# actions	Action lift (vs "No bid")	#imps	# actions	Action lift (vs "no bid")		
1	53,396	714	11.2%	59,703	826	28.7%	13.6%	156%
2	298,333	896	8.9%	431,637	980	19.1%	9.4%	115%
3	11,048,583	1,477	2.7%	11,483,360	1509	4.9%	2.2%	82%
4	3,915,792	2,016	6.6%	4,368,441	2,471	30.6%	22.6%	367%
5	6,015,322	6,708	19.6%	8,770,935	8,291	47.8%	23.6%	144%

Table: "Action lift" is the absolute # actions difference between lift-based bidding and value-based bidding. "Lift-over-lift" is comparing the their action lifts over "no bid".



Online A/B test (cont.)



DSPs get rewarded (paid) by advertisers based on the number of actions attributed to them. So there is a gap between DSPs and advertisers that must be abridged with attribution models.

Value-based bidding vs. lift-based bidding - DSP's perspective.

Adv	Val	Val ue-based b dding			ft-based bi	Inventory-	Cost-per-	
	#imps	# attrs	Inventory cost	#imps	# attrs	Inventory cost	cost diff	imp diff
1	53,396	50	\$278.73	59,703	50	\$300.31	7.7%	-3.6%
2	298,333	80	\$1,065.05	431,637	80	\$1,467.57	37.8%	-4.8%
3	11,048,583	240	\$25,522.22	11,483,360	240	\$25,837.56	1.2%	-2.6%
4	3,915,792	200	\$10,846.74	4,368,441	200	\$11,183.21	3.1%	-7.6%
5	6,015,322	500	\$19,296.51	8,770,935	500	\$23,501.90	21.8%	-16.5%

Table: Both bidders spent out equal amount of assigned budget, so the # attributions are always the same. Cost-per-impression is the inventory cost averaged by # impressions.



Online Advertising Landscape

Real-Time Bidding Mechanism

Case Study: AI Enables Better Bidding Strategies

Summary



RTB contributes a large portion of the online advertising revenue. It is the most dynamic marketplace that AI has to be leveraged.

Al technologies are the building blocks to implement advanced bidding strategies.

The effective impact of an ad can be modeled and predicted, with which a lift-based bidding strategy can bring advertisers true values.

Lift-based bidding may hurt DSPs with industry standard attribution model. The key to move DSPs to lift-based bidding is the attribution model.



Thank you! Questions?

Value-based bidding vs. lift-based bidding



Definition (AR, background AR, and AR Lift)

Given ad request, user, advertiser triplet (q, u, A), AR w.r.t. (q, u, A) is the probability that u will take the desired action defined by A after the ad of A is served to q, background AR w.r.t. (q, u, A) is the probability that u will take the desired action if the ad of A is not served to q, and AR lift as the difference between AR and background AR. We denote by p the AR, Δp the AR lift, and $p - \Delta p$ the background AR.

Definition (Value-Based Bidding)

Let p be the AR of a user if the advertiser's ad is shown, value-based bidding places a bid price of $a \times p$ to acquire an impression from this user for the advertiser, where a > 0.

Definition (Lift-Based Bidding)

Let Δp be the AR lift of a user if the advertiser's ad is shown, lift-based bidding places a bid price of $\beta \times \Delta p$ to acquire an impression from this user for the advertiser, where $\beta > 0$.

Bidder₁ is a value-based bidder, Bidder₂ is a lift-based bidder. What if they bid for the same advertiser simultaneously?

Value-based bidding vs. lift-based bidding (cont.)



Lemma

Bidder₁ wins the auction for u_i at the cost of $\beta \times \Delta p_i$ if $\alpha \times p_i > \beta \times \Delta p_i$; Bidder₂ wins the auction for u_i at the cost of $\alpha \times p_i$ if $\alpha \times p_i < \beta \times \Delta p_i$.

Theorem 1: Bidder₂ yields more actions than Bidder₁ when they are attributed same amount of credit.

i: the index of all the users

j: the index of those users that Bidder₁ wins (i.e., $a \times p_j > \beta \times \Delta p_j$)

k: the index of those users that Bidder₂ wins (i.e., $a \times p_k < \beta \times \Delta p_k$)

Expected attribution to Bidder₁: *j p_j*

Expected attribution to Bidder₂: k

Expected # of actions if only Bidder₁ is considered: $_{j}$ p_{j} + $_{k}$ $(p_{k} - \Delta p_{k})$

Expected # of actions if only Bidder₂ is considered: $(p_j - \Delta p_j) + (p_j - \Delta p_j) + (p_j - \Delta p_j)$

When the same amount of actions is attributed to Bidder_1 and Bidder_2 $\;$ (i.e.,

$$_{j}p_{j}=$$
 $_{k}p_{k})$, we have $\frac{_{j}p_{j}+_{k}(p_{k}-\Delta p_{k})}{_{j}p_{j}}<2-\frac{\alpha}{\beta^{k}}<\frac{_{j}(p_{j}-\Delta p_{j})+_{k}}{_{k}p_{k}}$.

Value-based bidding vs. lift-based bidding (cont.)



Theorem 2: Bidder₂ costs more than Bidder₁ when they are attributed same amount of credit.

- i: the index of all the users
- j: the index of those users that Bidder₁ wins (i.e., $a \times p_i > \beta \times \Delta p_i$)
- k: the index of those users that Bidder₂ wins (i.e., $a \times p_k < \beta \times \Delta p_k$)
- Expected attribution to Bidder₁: , p_j
- Expected attribution to Bidder₂: k
- Expected cost of Bidder₁: $_{j}$ $\beta \times \Delta p_{j}$ Expected cost of Bidder₂: $_{k}$ $a \times p_{k}$
- When the same amount of actions is attributed to Bidder₁ and Bidder₂ (i.e., $_{j}p_{j} = _{k}p_{k}$), we have $\frac{_{j}\beta \times \Delta p_{j}}{_{j}p_{i}} < \frac{_{j}\alpha \times p_{j}}{_{j}p_{i}} = \frac{_{k}\alpha \times p_{k}}{_{k}p_{k}}$.